University of Warsaw
Faculty of Mathematics, Informatics and Mechanics

Jacek Olczyk

Student no. 385896

Default values in Haskell record fields

Master's thesis
in COMPUTER SCIENCE

Supervisor:
PhD. Josef Svenningsson
Meta Platforms, Inc.

PhD. Marcin Benke
Instytut Informatyki

Warsaw, Dec 2022

Abstract

This thesis presents a prototype implementation of the syntax and semantics for specifying the default values for fields of Haskell record types. The proposed change aims to have no impact on existing Haskell code and minimal syntactic intrusion when the proposed feature is enabled. The use of default values happens only when using the braced record construction syntax already present in Haskell. The syntax for defining default field values follows established conventions from other programming languages in which this feature is present. The addition of this feature will noticeably improve the language's usability in large codebases.

Keywords

functional programming, Haskell, compilers, syntax, extension, default value, records, fields

Thesis domain (Socrates-Erasmus subject area codes)

11.3 Informatyka

Subject classification

D. Software
D.3 Programming Languages
D.3.3 Language Constructs and Features

Tytuł pracy w języku polskim

Domyślne wartości pól rekordów w Haskellu

Contents

Introduction

Algebraic data types have been a staple of functional programming languages since their earliest days. A typical extension of their utility for everyday programming is the record syntax, present in most popular functional programming languages. However, a notable functionality often missing from that feature is the ability to provide default values for individual record fields. While widely used workarounds are known, they suffer from various problems that a compiler-assisted solution can avoid.

This work aims to solve these problems in the case of the Haskell language [7] by proposing a custom syntax inside record definitions that enables the programmer to provide a default value to any field of a given record type. Chapter 1 contains the background and motivation behind these changes, as well as examples of problems they solve. For the sake of providing context to subsequent information, chapter 2 consists of a simple overview of the architecture of the GHC compiler. Next, the syntactic changes applied to the grammar of Haskell (Chapter 3), the semantic behavior of introduced features and the purely semantic changes to existing constructs (Chapter 4) are discussed. In all of these, we provide alternative considered solutions and the reasoning behind the final choices. We also explain the development process and decisions made when implementing the feature in GHC, the leading compiler of the Haskell language (Chapter 5), before ending on the conclusion 6;

Chapter 1

Background

1.1. Overview of record syntax and field defaulting

A major drawback algebraic data types usually suffer from when contrasted with data types commonly found in imperative languages (like `struct` in C) is their unwieldiness when it comes to data structures containing many fields. While most imperative and virtually all object-oriented languages feature a prominent and easy-to-use syntax for both naming and accessing individual fields of their data structures, the default for most functional languages has been to prefer data structures with unnamed fields. However, as projects grow larger and larger, having only unnamed data structure fields becomes increasingly cumbersome and error-prone. Thus, most functional languages provide an alternate record syntax for defining their data types.

In most cases, record syntax provides two important features: the ability to name individual fields of a data structure; and the automatic generation of field selector functions which absolves programmers from having to write the otherwise necessary, but cumbersome boilerplate code.

As an example, let's compare record syntax in C and Haskell representing a person's name and age.

```c
typedef struct {
  char* name;
  int age;
} Person;
```

```haskell
data Person
  = Person
  { name :: String
  , age :: Int
  }
```

In C, without this syntax, the only way to create compound data structures consisting of data of different types is manual pointer arithmetic:

```
// create a record with this syntax and set fields
Person person;
person.name = "Adam";
person.age = 21;

// one way to recreate the above on Linux without struct syntax
void* person = alloca(sizeof(char*) + sizeof(int));
*(char**)person = "Adam";
*(int*)(person + sizeof(char*)) = 21;
// the above would of course be extracted to accessor functions
// in hypothetical real-world code
```

In Haskell, the lack of named fields would necessitate boilerplate accessor functions or excessive pattern matching. For the above example, creation would be simple, but update would look like this:

```
adam :: Person
adam = Person "Adam" 21

-- manual pattern matching
updatePerson22, updatePersonJohn :: Person -> Person
updatePerson22 person =
  case person of
    Person name _ -> Person name 22

updatePersonJohn person =
  case person of
    Person _ age -> Person "John" age

-- accessor function
getName :: Person -> String
getName (Person name _) = name

getAge :: Person -> Int
getAge (Person _ age) = age

updatePerson22', updatePersonJohn' :: Person -> Person
updatePerson22' person = Person (getName person) 22
updatePersonJohn' person = Person "John" (getAge person)
```

With each additional field, the complexity of these updates becomes higher and higher. More examples can be found in [29].

1.1.1. History of records and field defaulting in other languages

The earliest examples of programming language syntax for record types date back to the 1950s and 60s with their introduction in COBOL in 1959 [14]. Popular contemporary languages like Algol 68 [28] and FORTRAN 77 [2] adopted their own syntax soon thereafter. They did not provide any facilities for providing default values for structure fields. Early object-oriented languages either used constructors for this purpose (Smalltalk [6]), or nothing at all (Simula [3]). The same followed for most of the languages developed in the 1970s and 80s: languages with either no default initialization of record types, like C and Pascal, or default initialization via constructor functions, like in C++. A notable exception is Common Lisp, which being released in 1985 can be counted as one of the earliest examples of records with dedicated

syntax for default field values [17]. With the rise of popularity of programming and the coming of the Internet era, syntactic support for default values has suddenly changed from obscure to ubiquitous, as will be clear from the examples in the next section.

1.1.2. Modern examples in other languages

Examining 10 of the most popular general-purpose programming languages in 2022, most of them provide in their latest version some utility for specifying default field values, either through constructor functions or dedicated syntax. Chosen programming languages are: C++, Java, JavaScript, TypeScript, Python, C#, PHP, Go, Kotlin and Rust. The choice of languages was based on statistics from [16], after removing markup languages, query languages and scripting languages, as well as C which was already discussed above.

Of course, any programming language offers support for default values through creating an object and using it as a "default" object. This approach is not investigated, as it requires no help from the language to implement, unless there is a language-wide convention on the details on how the object is supposed to be created. Furthermore, this approach has other drawbacks discussed in section 1.2.2.

The below table illustrates support for field defaulting in the above languages.

	Dedicated syntax	Parameterless constructors
C++	Since C++11	Yes
Java	Yes	Yes
JavaScript	Since ES6	Since ES6, or as prototypes
TypeScript	Yes	Yes
Python	dataclasses only	Yes
C#	Yes	Yes
PHP	Since PHP5	Since PHP5
Go	No	No
Kotlin	Yes	Yes
Rust	No	Limited, through the Default trait

Table 1.1: Support for specifying default field values across popular languages.

Below are examples of each syntactic construct mentioned in the table. Note that, of course, languages supporting the parameterless constructor syntax allow not only for providing default values for all fields at once, but also for providing values of an arbitrary subset of fields through simply adding parameters to the constructor.

- C++: dedicated syntax (C++11 onwards) [9]

```
struct Person {
  std::string name = "Adam";
  int age = 21;
};
```

  The above values get assigned to fields if the fields are missing in the constructor's member initialization list (showcased in the next bullet).

- C++: default constructor [8]

```
struct Person {
```

```cpp
    std::string name;
    int age;
    Person() : name("Adam"), age(21) {}
};
```

This approach uses the constructor's member initialization list, which guarantees that
the fields are initialized with the default values before entering the constructor body.

- Java: dedicated syntax

```java
class Person {
  public String name = "Adam";
  public int age = 21;
}
```

This approach assures that the default values are assigned before entering the body of
any user-defined constructor.

- Java: parameterless constructor

```java
class Person {
  public String name;
  public int age;
  public Person() {
    name = "Adam";
    age = 21;
  }
}
```

This approach is just manual assignment of values, the only syntactic help is the con-
structor itself.

- JavaScript: dedicated syntax (since ES6)

```javascript
class Person {
  name = "Adam";
  age = 21;
}
```

- JavaScript: parameterless constructor (since ES6)

```javascript
class Person {
  constructor() {
    this.name = "Adam";
    this.age = 21;
  }
}
```

- JavaScript: prototypes

```javascript
function Person() {
  this.name = "Adam";
  this.age = 21;
}
```

This defines a constructor function for `Person`, with any required methods being later added to the function's prototype.

- TypeScript: dedicated syntax

```typescript
class Person {
  name = "Adam";
  age = 21;
}
```

Note that the type of the fields is automatically deduced from the initializing expression.

- TypeScript: parameterless constructor

```typescript
class Person {
  name: string;
  age: number;
  constructor() {
    this.name = "Adam";
    this.age = 21;
  }
}
```

This approach is just manual assignment of values, the only syntactic help is the constructor itself.

- Python: dataclasses

```python
@dataclass
class Person:
  name: str = "Adam"
  age: int = 21
```

- Python: parameterless constructor

```python
class Person:
  def __init__(self):
    self.name = "Adam"
    self.age = 21
```

- C#: dedicated syntax

```csharp
class Person {
  public String Name = "Adam";
  public int Age = 21;
}
```

This approach assures that the default values are assigned before entering the body of any user-defined constructor.

- C#: parameterless constructor

```csharp
class Person {
  public String Name;
  public int Age;
```

```
      public Person() {
        Name = "Adam";
        Age = 21;
      }
    }
```

This approach is just manual assignment of values, the only syntactic help is the constructor itself.

- PHP: dedicated syntax (since PHP 5)

```
    class Person {
      public $name = "Adam";
      public $age = 21;
    }
```

- PHP: parameterless constructor (since PHP 5)

```
    class Person {
      public $name;
      public $age;
      function __construct() {
        $this->name = "Adam";
        $this->age = 21;
      }
    }
```

Note that this is not actually custom syntax, constructors in PHP are regular methods with the name __construct.

- Kotlin: dedicated syntax and parameterless constructor

```
    class Person {
      val name = "Adam";
      val age = 21;
    }
```

The body of the class is explicitly considered to be the primary constructor in Kotlin.

- Kotlin: parameterless (secondary) constructor

```
    class Person {
      val name: String;
      val age: Int;
      constructor() {
        name = "Adam";
        age = 21;
      }
    }
```

- Rust: Default trait

```
    struct Person {
      name: String;
      age: i32;
```

```
}

impl Default for Person {
  fn default() -> Person {
    Person {
      name: "Adam",
      age: 21,
    }
  }
}
```

This approach is limited by only being able to provide default values to all fields of the struct at once. It also is not an example of syntactic support, as all that is required is a trait. Because of that, use of this trait for default construction relies entirely on convention.

1.1.3. Conclusion

From the above considerations it is clear that the vast majority of modern languages offer a convenient syntax for providing default field values. Comparing that to the historical situation where it was not offered by any major language, it is clear that it has become a quality-of-life feature that today's programmers have come to expect from their language of choice. Thus, an addition of such a feature is likely to make teams and developers more eager to adopt Haskell as their language of choice.

1.2. Field defaulting in Haskell

1.2.1. Introduction

Despite record syntax's relative ubiquity in modern functional programming languages, most implementations of it (e.g. OCaml, F#, ReasonML and of course Haskell) have an important usability feature missing when compared to their imperative counterparts — the ability to leave out certain fields empty when creating a concrete object of a given type. Especially when dealing with data structures containing a great number of fields, e.g. types that represent a configuration file, the necessity of providing every field of the structure with a value can become unwieldy very quickly. The only example of this syntax in a popular functional programming language seems to be Erlang, where records can be default-constructed. Before version Erlang/OTP 19, the fields without default values were provided with a value of the singleton type 'undefined' [5]. Since Erlang/OTP 19, fields without a default value must be given a value upon construction.

1.2.2. Prior attempts at solving the issue

This issue has been heretofore partially mitigated thanks to another common feature of record syntax in programming languages: the record update syntax. It allows for the construction of a record in such a way that any unspecified fields are copied over from an already existing record. If the author of a type then provides their users with a globally visible 'default' object with all its fields already filled out, a user of the type can utilize the record update syntax to mimic the functionality offered by default field values.

However, this approach is not without its drawbacks. Firstly, it can prove a challenge if some fields don't lend themselves easily to default values. For example, a data type might have

an ID component along with invariants guaranteeing its uniqueness. In such a situation, great care must be taken when creating a new record to ensure that the ID component is always modified from the value provided by default. It is also possible to circumvent that using optional types (like Haskell's `Maybe`), but this makes the field value unnecessarily cumbersome to access after the creation of the record. Unlike its imperative counterpart, this approach does not offer the programmer any automated systems that would detect the absence of a value that is necessary for the type. Given that a major advantage functional languages tout over others is the strength of their automated compile-time error detection, this is an area in which clear improvement can be made.

In Haskell, as well as other languages, it is also possible to create a record without specifying all of its fields. This way it is not necessary to provide the troublesome fields with any value, relying on the user to add it later. However, in that case, the fields are still automatically filled with bottom values. To make matters worse, such creation creates compile-time warnings (with `-Wmissing-fields` enabled) when defining the default record, and no compile-time warnings when some error values are not replaced in a record update. Because of that, this solution is even worse than the previous one.

This issue is completely solved with the advent of default field value syntax: all fields that can't be easily defaulted, are just not given a default value, and with `-Wmissing-fields` we get compile-time warnings whenever any non-defaulted field is missing.

Secondly, the necessity of creating a custom object to use as a default can lend itself to improper coding style and more difficulty in codebase navigation. In order to create a default object, one must first decide on a name that it will be given, which in large codebases necessitates either coming up with a naming convention or dealing with the repercussions of inconsistent naming. Furthermore, there is no requirement for the programmers to place the default object in the same close location, or even in the same file as the definition of the type itself. This can quickly make the issue of finding the default object a true detective's task.

By contrast, none of these problems are present when using default field values.

1.2.3. Direct use cases

Aside from the obvious functionality improvements, allowing default fields in records has positive impact on the overall Haskell ecosystem. Many large multi-language codebases, such as the giant monorepos used by large companies, use interface definition languages (IDLs) like the ones used in Apache Thrift or Google's Protobuf for interfacing between systems written in different languages. Such IDLs usually provide this functionality through a language-agnostic syntax for declaring the general layout of data structures and/or procedures. This syntax usually (as is the case in Apache Thrift) contains a way to provide default values for fields of the data structures being defined. Because of that, any language that doesn't natively support default field values in data structures will need to implement workarounds before codebases written in that language are able to intercommunicate via IDLs.

Code migrations

A major potential real-world use case is helping with code migrations. As an example, consider a system with two codebases, *user* and *infra*. The *infra* code provides type definitions for configs, like this simplified example:

```
data Config
  = Config
  { name :: String
  , value :: SomeType
  }
```

Code in *user* will then construct values of this type, like so:

```
someVar :: Config
someVar = Config { name = "my config", value = myValue }
```

To ensure runtime safety of the system, constructing this value is protected via the combination of the `-Wmissing-fields` and `-Werror` compiler flags. In their presence, omitting a field creates a compile-time error, making sure that every field value is correctly in place. Since it is easy for a programmer to forget about a field, especially when there are many of them, this check measurably increases correctness of the system. However, when there is a need to add a new field, problems with that approach arise. Even if the system was limited to just one codebase, a sufficiently large system cannot be easily migrated to use a new value in all the occurrences of the construction. With multiple codebases, the process complicates even more. Here's a step-by-step example of how that could look like, if one wanted to add a field `owner` of type `String`:

1. We cannot add the field directly without creating compile errors in *user* code. To work around that, we create a default value called, let's say, `defaultConfig` in the same module of *infra* code as the definition of the type.

   ```
   defaultConfig :: Config
   defaultConfig = Config { name = "default name", value = defaultSomeTypeValue }
   ```

2. This code needs to be released in order for `defaultConfig` to be accessible in *user* code.

3. Then, all constructors of `Config` in *user* can be modified to a record update of `defaultConfig`:

   ```
   someVar :: Config
   someVar = defaultConfig { name = "my config", value = myValue }
   ```

   This change silences the warning

4. This code, again, needs to be released before any further changes in *infra* are made.

5. With that setup, we can safely introduce a new field in *infra*, immediately modifying `defaultConfig` with a default value :

```
data Config
  = Config
  { name :: String
  , value :: SomeType
  , owner :: String
  }

defaultConfig :: Config
defaultConfig
  = Config
  { name = "default name"
  , value = defaultSomeTypeValue
  , owner = "default owner"
  }
```

6. Yet again, a release is needed.

7. Now, all *user* code can revert to using constructors, as soon as the relevant code is adapted to make use of the new field.

```
someVar :: Config
someVar
  = Config
  { name = "my config"
  , value = myValue
  , owner = "my owner"
  }
```

8. This *user* code must again be released. This restores the protection given by `-Wmissing-fields`.

By contrast, with the addition of field default value syntax the above process could be shortened to just one *infra* release containing the new field with the default value:

```
data Config
  = Config
  { name :: String
  , value :: SomeType
  , owner = "default owner" :: String
  }
```

This way, `-Wmissing-fields` is working for all the other fields, and teams working on *user* code can add support for the new field at their own pace. Once this is done, the default value can be removed to ensure it's also protected by `-Wmissing-fields`.

Example of a big config

A good example of a large record type in a widely used Haskell codebase is the DynFlags structure in GHC. Its type definition is indeed massive, and so is the function used for constructing its default value. The definition is included in Appendix A [25], and the default value construction in Appendix B [24]. This is because the former takes up 6 pages and the latter 4 pages. Investigating the default value for the type, a couple of notable observations can be made.

- It's a function taking a `Settings` value as an argument, as some fields of the default value rely on settings.

- Most of the other field values are either constants or global values.

- Three fields contain bottom values that panic on evaluation.

Clearly, the values that do not depend on the setting value are great candidates for being defaulted in the type definition itself. Only 8 fields actually depend on the setting value, so with the default value extension the function can be reduced to:

```
defaultDynFlags :: Settings -> DynFlags
defaultDynFlags mySettings =
-- See Note [Updating flag description in the User's Guide]
  DynFlags {
    backend = platformDefaultBackend (sTargetPlatform mySettings),
    ghcNameVersion = sGhcNameVersion mySettings,
    fileSettings = sFileSettings mySettings,
    toolSettings = sToolSettings mySettings,
    targetPlatform = sTargetPlatform mySettings,
    platformMisc = sPlatformMisc mySettings,
    rawSettings = sRawSettings mySettings,
    generalFlags = EnumSet.fromList (defaultFlags mySettings),
  }
```

The fields with bottom values illustrates another possible use case: explicitly deferring missing fields check to runtime. While generally it's considered beneficial to catch missing fields during compile-time, it's clear that in some cases this cannot be easily done because of technical debt. Sometimes it's enough to wrap the field inside a `Maybe`, but it is not without its downsides. Providing a panic value as a temporary placeholder allows for better error messages in case of a programming error, as well as silencing any `-Wmissing-fields` warnings. With default field values this can be done once, globally, instead of having to manually ensure it in every constructor.

Chapter 2

GHC architecture

This section contains an abridged overview of the architecture of the frontend of the GHC compiler to serve as background for implementation details provided later. The compilation process in the GHC compiler consists of 6 main phases [4], the first 4 of which are:

1. Parsing

2. Renaming

3. Typechecking

4. Desugaring

These are all the phases that deal with the Haskell AST, mostly unchanged from the parser. The last one, the desugarer, desugars the Haskell AST to the much smaller internal language Core, which is used for optimization. Before discussing the passes themselves, it's worth to describe how GHC deals with modification of the AST throughout the frontend of the compiler.

2.1. Trees that grow

Throughout the phases of the frontend, various data about the Haskell AST is generated (and other becomes no longer needed). In order to avoid data structure fields or constructors being unused during some phases or having multiple slightly different versions of the same AST type definitions, a mechanism was introduced in [11] which parameterized the entire syntax tree with a type argument representing the current phase of compilation. Three distinct phases are available: `GhcPs`, `GhcRn`, `GhcTc`, with the parser producing code parameterized by `GhcPs`, the renamer converting `GhcPs` to `GhcRn`, and the typechecker converting `GhcRn` to `GhcTc`. With that in mind, when creating a data structure, the developer writes all constructors and fields common to all phases as they normally would, with a minor exception. Every type `T` gets an extra constructor `XT` (eXtension to `T`), which allows for phase-specific constructors, and every constructor `CtorT` has as its first field set to type `XCtorT`, to allow for phase-specific constructor fields. Any type written in extension field or an extension constructor definition is actually a `type family`, which when parameterized by the phase, allows for a different `type instance` for each phase.

2.2. Compilation phases

2.2.1. Parsing

GHC implements parsing of source code through a monolithic Happy grammar [1], although this wasn't initially the case [10]. The data structures emitted by the parser are often ambiguous as to which production was actually used to parse a given part of the source code. This is a result of a policy of 'overparse, then filter out the bad cases' [20], which is also used by this implementation, see section 5.2. For example, in certain contexts both patterns and expressions are valid, so knowing which grammar derivation to use would require infinite lookahead. This is solved through a validation monad which allows parsing contexts to specify which constructs are possible, failing the parse on the rest. Another system of note are the source location and annotation types, which are ingrained into the types that form the Haskell source code AST. Any parser node has to extract the location and annotation data from its subnodes, and very carefully combine them together, as there is no automatic mechanism for detecting errors in code like this.

2.2.2. Renaming

The task of the renamer is to match each name usage in the code with the matching definition. Its code is intertwined with the typechecker code, and uses the same monad [21].

2.2.3. Typechecking

Typechecking a module happens concurrently (but not in parallel) with renaming, because of Template Haskell splices [18], which need to be processed together in connected components [22]. The typechecker uses the `Type` type, which is used to type Core, instead of the `HsType` created by the parser [22]. This is because the typechecking process is quite complicated, and thus requires a data structure that is easy to manipulate.

2.2.4. Desugaring

The desugarer transforms the Haskell AST (with the `GhcTc` phase indicator) into an explicitly typed variant of System FC called Core [19]. The Core language AST consists of only 10 constructors for its expressions [23], and great care is taken to not modify this language when not necessary, so any changes that are little more than syntactic sugar should be implemented as part of the desugarer.

Chapter 3

Syntax

Among a wide variety of languages supporting providing default values to fields of data structures, there has been a virtually unanimous agreement on the syntax used for providing the default values. Any language that supports this feature on the compiler level seems to agree on the base syntax being `field_name = value`, with minor variations depending on the given language's syntax for field type annotations. Thus, declaring an integer field with default value of 1 in C++, Java, and many others, looks like this: `int x = 1;`. As another example, the same is accomplished in TypeScript with the syntax `x: number = 1,`. Many other examples presented in 1.1.2 also follow this pattern.

3.1. Syntax for declaring a default field value

With that in mind, we needed to accommodate the `field_name = value` syntax to existing syntax for Haskell records. A single field in a Haskell record currently looks like the following:

```
fieldName :: FieldType
```

Considering the above examples, a few alternative approaches for the default value syntax come to mind.

- type after value

```
fieldName = field_value :: FieldType
```

- type after name

```
fieldName :: FieldType = field_value
```

- variants of both of the above, but allowing for omitting the type of the expression

In total four variants, and the one chosen here is the "type after value" approach without omitting the type signature.

3.1.1. Order of syntactic components

To decide between 'type after value' and 'value after type', we looked at language constructs already prevalent in existing Haskell code.

A regular top-level Haskell binding with a type annotation can have multiple forms:

1. standalone type annotation

```
name :: Type
name = value
```

2. type after value

```
name = value :: Type
```

3. value after type

```
name :: Type = value
```

By far the most popular of these variants is 1, but it's not readily adaptable to record fields. To choose between 2 and 3, we note that 3 is very rarely used in real code. Furthermore, at first glance, it might seem as if the default value is being assigned to the type, not the field name. In contrast, 2 looks like a familiar construct — a binding name without a type annotation, and its right-hand-side expression with one (and it is already parsed as such by GHC).

3.1.2. Omitting the type signature

It would seem like a good idea to allow users to omit the type signature for a field with a user-supplied default value:

```
data NewRecordType = NewRecordTypeCtor { fieldName = expression }
```

After all, it seems possible to always infer the type from the defaulting expression. As it turns out, the reality is not that simple. Many tools that are part of the GHC project currently need to be able to present the type of all fields of a given datatype without the information provided by a full typecheck of the program. The list of these tools includes Haddock, but also the typechecker itself. While the typechecker itself could theoretically be modified to avoid this problem, forcing a typecheck on other tools (such as Haddock) would unnecessarily slow down their performance without a noticeable gain. As providing type annotations is a good practice in almost all use cases and virtually all existing code already has to do so for all their record fields, this potential loss in functionality seems to have negligible cost.

3.1.3. Multiple names in one line

Regular Haskell syntax for field declarations allows for multiple fields of the same type to be written on one line:

```
data NewRecordType
  = NewRecordTypeCtor {
    fieldName1, fieldName2, fieldName3 :: FieldType
  }
```

It's not a widely used feature, especially given its counterpart in C/C++ is considered bad style. Its use is not particularly readable either: programmers usually visually judge the number of fields in a type by the number of lines. It also makes the code harder to document as it's not clear how individual fields should be documented, and the Haddock docs [15] do not take that syntax into account. With that in mind, complicating this syntax even further with default values does not seem desirable:

```
data NewRecordType
  = NewRecordTypeCtor {
    fieldName1 = expr1, fieldName2 = expr2, fieldName3 = expr3 :: FieldType
  }
```

Furthermore, it could create ambiguity for the reader when only the last field is defaulted:

```
data NewRecordType
  = NewRecordTypeCtor {
    fieldName1, fieldName2, fieldName3 = expr :: FieldType
  }
```

In this example, it is not immediately clear whether `expr` is the default value for all fields, or only `fieldName3`. For these reasons, this feature was not included in the proposed syntax.

3.2. Syntax for record construction

The chosen method for initializing fields with their default value is through braced record syntax. Because the only other way to construct a record is through a regular constructor function, which cannot have optional arguments (just like all other Haskell functions), braced record syntax is the sole method of assigning default values to fields. Omitting a given defaulted field is the only requirement for activating the mechanism. Since omitting fields in record construction was already allowed (but produced a compile-time warning), no syntax change is necessary.

3.3. Summary

All in all, the new syntax is as follows: inside a record constructor definitions, any field declared in the form `name :: Type` can optionally be annotated with a default value: `name = <expr> :: Type`, where `<expr>` is an arbitrary expression (of the same type as the field). Below is a complete example of a Haskell record type with one of its fields defaulted.

In our view, the final syntax combines the best compromise between readability, new user learning curve, and implementation viability.

```
data NewRecordType
  = NewRecordTypeCtor
  { fieldName :: SomeType
  , fieldName2 = expression2 :: SomeType2
  }
```

No syntax change is required for inserting default values during record construction.

Chapter 4

Semantics

4.1. Examples

The semantics for the proposed extension of the language are very simple. Recall the (slightly renamed) example type from the previous chapter:

```haskell
data Record1
  = RecordCtor2
  { field1 :: Int
  , field2 = expression1 :: String
  }
```

In order to utilize the mechanism of inserting default values, usage of the record construction syntax is necessary:

```haskell
exampleRecord :: Record1
exampleRecord
  = RecordCtor1
  { field1 = 12
  -- field2 is initialized with the default value
  }
```

Heretofore, omitting a field when constructing a record resulted in that field containing a lazy error value, which terminates the program on evaluation. The above example is the only possible way of constructing a record with a non-error default value. Below we give non-examples of constructions that do not change their semantics with this change (but could reasonably be expected to). In the first example, the field is initialized with an error value, both in the current version of Haskell and with the changes described in this thesis applied to GHC.

```haskell
data Record2
  = RecordCtor2
  { field1 :: Int
  , field2 :: String
  }

exampleRecord :: Record2
exampleRecord
  = RecordCtor2
  { field1 = 12
  -- field2 is initialized with an error value, raising an error when evaluated
```

```
    }
```

In the next example, we note that record update is not affected by these changes. Every unspecified field in the record update will be copied directly from the old record, even if its value is an error value. The example record created by the following code is equal to `RecordCtor3 { field1 = 5, field2 = 4 }`.

```
data Record3
  = RecordCtor3
  { field1 = 1 :: Int
  , field2 = 2 :: String
  }

baseRecord :: Record3
baseRecord = RecordCtor3 { field1 = 3, field2 = 4 }

exampleRecord :: Record3
exampleRecord
  = baseRecord
  { field1 = 5
  -- field2 is initialized using the value in baseRecord, not the default value
  }
```

4.1.1. The -Wmissing-fields flag

Furthermore, this change has a side effect on the behavior of programs with the compiler flag `-Wmissing-fields` on. With that flag, any record construction expression produces a compile-time warning when any field initialization is omitted. Given that any field for which a default value has been provided cannot be missing, this means that no warning will be generated in these cases. While just opting in to the extension does not change the behavior by itself, providing a default value for a field has the effect of silencing any warnings that could have previously been reported by the compilers for constructors of a given type.

4.2. A more formal statement

For precision, we offer a reasonably rigorous semantic description of field defaulting:

Description 1. Let T be a type with a record constructor T' containing a field f for which a default expression e has been provided. Then, if while calling the constructor T' via the record construction syntax, field f is not given any value, during evaluation the construction happens *as-if* field f had been specified and given the value e. The presence of `-Wmissing-fields` flag does not result in any warnings for any usage of f.

4.3. Operational semantics

Given Haskell's usual semantics as a pure and lazy language, the defaulting expression is only evaluated at most once, if it's needed. This allows for arbitrarily complex expressions without risking unnecessarily long runtime of the program. Even if the defaulted field's type is a computation inside the `IO` monad (and thus can be considered to have side effects), the evaluation of the expression only produces instructions for the runtime system without

actually executing them. Thanks to that we can be certain that the default values will behave as expected.

This operational, rather than denotational semantic aspect of program behavior is implemented through storing default values as global bindings. That way, it implicitly performs the optimization described in [12] as the full laziness transformation. This naturally is expected to result in the same consequences and considerations as outlined therein:

- Programs can be made faster by not repeating work (to compute the value) and allocations (to store the value).

- When the default value 'is already a value, or reduces to a value with a negligible amount of work', the only potential gain is in allocations.

- The values can potentially cause a space leak (e.g. when a default value is an infinite list and more elements get evaluated than will be needed in the future).

- They are harder to garbage-collect.

How do these apply to the typical use case of default field values? In a reasonable program, one would expect default values to be more likely to need to be reused. Thus, making sure that their value isn't recomputed on each evaluation of a constructor that does not contain the field seems like a good approach. It also follows the principle of least confusion [13]: it's not immediately obvious to a user writing a constructor with a defaulted field how much work will be performed on construction, so the best approach is to minimize it.

Furthermore, it's common for default values to require little to no computation. For example, a major use case of default values is placeholders meant to be filled out later. Often programmers choose the simplest possible value for a given type as the default, e.g. `0 :: Int`. With this in mind, it seems unlikely for programmers to provide default values that could cause a noticeable space leak, and most of the time one might reasonably expect a bunch of allocations to be omitted. Unfortunately, there is no way to verify these assumptions without developing an alternative implementation that doesn't make the default values global.

Chapter 5

Implementation

We implemented the proposed change to the language inside the leading Haskell compiler, GHC (Glorious Haskell Compiler) [7]. The algorithm is a journey that takes the default expression from being declared to being substituted when a field is missing. It consists of the following main steps:

1. Parse the defaulting expression inside the constructor.

2. Give it a unique name and group it together with top-level bindings so that dependency analysis, renaming and typechecking are performed without disturbances.

3. Mark defaulted fields as not actually missing.

4. During desugaring to Core, replace any missing fields with their default values.

The necessary changes made to the GHC code can be broken down into X parts:

1. Changes to datatypes

2. Parser

3. Renamer

4. Desugarer

5.1. Changes to datatypes

For the initial prototype we tried to keep the changes to existing internal GHC data structures to a minimum.

ConDeclField

Taken from [27].

```
-- | Constructor Declaration Field
data ConDeclField pass -- Record fields have Haddock docs on them
  = ConDeclField { cd_fld_ext :: XConDeclField pass,
                   cd_fld_names :: [LFieldOcc pass],
                             -- ^ See Note [ConDeclField passs]
                   cd_fld_type :: LBangType pass,
                   cd_fld_doc :: Maybe LHsDocString }
```

First necessary change was to the `ConDeclField` type, which represents a single field of a record constructor. We modify it by adding a field called `cd_fld_ini` of type `Maybe (LHsExpr pass)`:

```haskell
-- | Constructor Declaration Field
data ConDeclField pass -- Record fields have Haddock docs on them
  = ConDeclField { cd_fld_ext :: XConDeclField pass,
                   cd_fld_names :: [LFieldOcc pass],
                                 -- ^ See Note [ConDeclField passs]
                   cd_fld_type :: LBangType pass,
                   cd_fld_doc :: Maybe (LHsDoc pass),
                   cd_fld_ini :: Maybe (LHsExpr pass) }
```

The `L` prefix in `LHsExpr` means that the expression contains information about its location in the source file. The `pass` parameter signifies the current phase of the compilation (using the Trees that grow technique described in 2.1 and in detail in [11]). This modification allows us to pass the defaulting expression from the parser to the renamer.

FieldLabel

Taken from [26].

```haskell
-- | Fields in an algebraic record type; see Note [FieldLabel].
data FieldLabel = FieldLabel {
    flLabel :: FieldLabelString,
    -- ^ User-visible label of the field
    flHasDuplicateRecordFields :: DuplicateRecordFields,
    -- ^ Was @DuplicateRecordFields@ on in the defining module for this datatype?
    flHasFieldSelector :: FieldSelectors,
    -- ^ Was @FieldSelectors@ enabled in the defining module for this datatype?
    -- See Note [NoFieldSelectors] in GHC.Rename.Env
    flSelector :: Name
    -- ^ Record selector function
  }
  deriving (Data, Eq)
```

The renamer moves information about types and their constructors from the data structures that represent abstract syntax to internal structures that are easier to work with. There, information about record fields is stored inside the `FieldLabel` type, to which we added a new field, `flIniExpr :: Maybe Name`.

```haskell
-- | Fields in an algebraic record type; see Note [FieldLabel].
data FieldLabel = FieldLabel {
    flLabel :: FieldLabelString,
    -- ^ User-visible label of the field
    flHasDuplicateRecordFields :: DuplicateRecordFields,
    -- ^ Was @DuplicateRecordFields@ on in the defining module for this datatype?
    flHasFieldSelector :: FieldSelectors,
    -- ^ Was @FieldSelectors@ enabled in the defining module for this datatype?
    -- See Note [NoFieldSelectors] in GHC.Rename.Env
    flSelector :: Name,
    -- ^ Record selector function
    flIniExpr :: Maybe Name
    -- ^ Initializing expression
  }
  deriving (Data, Eq)
```

The `Name` refers to the generated name of the binding generated for the expression.

5.2. Parser

Parser uses Happy, a grammar-based parser generator for Haskell. The existing grammar node for record fields is called `fielddecl`. Its only production is `fielddecl : sig_vars '::' ctype`. The above parses a comma-separated list of variable names, followed by the :: operator, followed by the type of the field(s). The names are a list to allow for declaring multiple fields with the same type in one line, e.g. `a, b, c :: Foo`. We want to modify it to allow for `= expr` between the name of the field and the type. One obvious solution would be to add another production producing `var '=' exp '::' ctype` (as we only want to allow a single name to be defaulted on one line). This unfortunately produces a reduce/reduce conflict, as `exp` can already contain a type signature. We solve this conflict by extracting the production for an expression with a type signature to a separate node, `typedexp`:

```
typedexp :: { ECP }
        : infixexp '::' ctype
            { ECP $
                unECP $1 >>= \ $1 ->
                rejectPragmaPV $1 >>
                mkHsTySigPV (noAnnSrcSpan $ comb2Al $1 (reLoc $>)) $1 $3
                        [(mu AnnDcolon $2)] }
```

This allows us to parse `var '=' typedexp`, which forces a type signature and guarantees no conflicts. The aforementioned production is introduced as a new node, `varini`, for clarity:

```
varini :: {(LocatedN RdrName, LHsExpr GhcPs, LHsType GhcPs)}
  : var '=' typedexp
    {% runPV $ (unECP $3 :: PV (LHsExpr GhcPs)) >>=
            \ e@(L _ (ExprWithTySig _ _ (HsWC _ (L _ t)))) ->
                return ($1, e, sig_body t) }
```

It is used in `sig_vars` definitions like this:

```
fielddecl :: { LConDeclField GhcPs }
  -- A list because of f,g :: Int
  : sig_vars '::' ctype {- code for the existing prodution -}
  | varIni
    {% case $1 of
        (ln@(L l n), expr, t) -> acsA (\cs -> L (locA l)
                (ConDeclField (EpAnn (glNR ln) [] cs)
                [L (l2l l) $ FieldOcc noExtField ln] t Nothing (Just expr)))}
```

There is, however, another issue that needs resolving: the `exp` production, and by extension also `typedexp` production both return a value of type `ECP` instead of `HsExpr`. This is due to ambiguity concerns mentioned in section 2.2.1, so the `unECP` function must be used, annotating the result with type `PV (LHsExpr GhcPs)` to mark the parsing context as requiring an expression.

5.3. Renamer

5.3.1. Preprocessing

Before we start the renaming process, we want to make sure we are making as much use of existing GHC piping as possible. Thus, at the beginning of the renamer, we extract all default values for fields into separate bindings. In the type declarations, the user-provided expressions are then replaced (inside the `ConDeclField` structures) with simple variable expressions. The user-provided expressions are given freshly generated names and grouped together with other value declarations for the module. All the declarations are fed to the renamer in this way, and there to the next phase of renaming proceeds.

5.3.2. Conversion

Early in the renaming process, the abstract syntax for type declarations is parsed into more convenient data structures. It is here that record selector functions are created, and their names are placed inside the `FieldLabel` objects created alongside them. This means that it's the perfect place for inserting the names of the default expressions into the `FieldLabel`s. To do that, before checking individual constructors or fields, a mapping is constructed of all field names from a given type declaration to the names of their default expressions.

5.3.3. Missing fields detection

The fields that have been given default values should be excluded from both the list of missing strict fields (which produce an error when not specified) and the list of missing non-strict fields (which only produces a warning). When detecting if fields are missing, the algorithm compares the available `FieldLabel`s for the given constructor with the occurrences provided in a given expression. Thus, to exclude defaulted fields, it is sufficient to filter out the `FieldLabel`s that have their `flIniExpr` field set to a `Just` value.

5.4. Desugarer

During desugaring, we need to fill any unspecified record fields with the variable name corresponding to the appropriate default expression. Normally, the desugarer fills all missing fields with error values. Given that the desugarer has access to the `FieldLabel`s of the arguments being desugared, we can check if they contain the default expression and desugar the field to contain the variable name instead of an error.

Chapter 6

Conclusion

This work proposed a custom syntax for default record fields in Haskell, detailed a proof-of-concept implementation and discussed the design choices and trade-offs behind the syntax and semantics of the implementation.

From this work it is clear that an implementation of default record field values in Haskell is possible, usable, useful and hopefully soon implemented into the mainline GHC. The syntax is clear and intuitive and the behavior of programs that use it is predictable and efficient. Programs using the extension are simpler, more expressive and have the potential to be faster than ones written without the extension. Being able to guarantee that a value is always present makes systems built with this feature more reliable and less error-prone. Usage of the feature will also be helpful in workflows related to maintenance of larger codebases.

The next step for adoption of the feature is to submit a proposal for the syntax to be accepted as a language extension in the GHC compiler, to bring this feature to the widest possible audience. Further discussion with the GHC maintainers will be required to finalize the design choices and implementation details, and this work is bound to serve as a good starting point for that discussion.

In conclusion, this proof-of-concept implementation of default values for record fields in Haskell demonstrates the feasibility and potential benefits of this feature. By allowing users to specify default values for record fields, the implementation reduces the amount of boilerplate code that is required and improves the readability and maintainability of Haskell programs. Additionally, the implementation is consistent with the existing design of Haskell language constructs, and it can be integrated seamlessly into the existing Haskell codebases. Overall, this implementation of default values for record fields in Haskell is a practical approach that could significantly improve the usability and expressiveness of the language.

Appendix A

GHC DynFlags type definition

Sourced from [25].

```haskell
data DynFlags = DynFlags {
  ghcMode               :: GhcMode,
  ghcLink               :: GhcLink,
  backend               :: !Backend,
    -- ^ The backend to use (if any).
    --
    -- Whenever you change the backend, also make sure to set 'ghcLink' to
    -- something sensible.
    --
    -- 'NoBackend' can be used to avoid generating any output, however, note that:
    --
    -- * If a program uses Template Haskell the typechecker may need to run code
    --    from an imported module. To facilitate this, code generation is enabled
    --    for modules imported by modules that use template haskell, using the
    --    default backend for the platform.
    --    See Note [-fno-code mode].


  -- formerly Settings
  ghcNameVersion  :: {-# UNPACK #-} !GhcNameVersion,
  fileSettings    :: {-# UNPACK #-} !FileSettings,
  targetPlatform  :: Platform,      -- Filled in by SysTools
  toolSettings    :: {-# UNPACK #-} !ToolSettings,
  platformMisc    :: {-# UNPACK #-} !PlatformMisc,
  rawSettings     :: [(String, String)],
  tmpDir          :: TempDir,

  llvmOptLevel        :: Int,        -- ^ LLVM optimisation level
  verbosity           :: Int,        -- ^ Verbosity level: see Note [Verbosity levels]
  debugLevel          :: Int,        -- ^ How much debug information to produce
  simplPhases         :: Int,        -- ^ Number of simplifier phases
  maxSimplIterations  :: Int,        -- ^ Max simplifier iterations
  ruleCheck           :: Maybe String,
  strictnessBefore    :: [Int],      -- ^ Additional demand analysis

  parMakeCount        :: Maybe Int,  -- ^ The number of modules to compile in parallel
                                     --   in --make mode, where Nothing ==> compile as
                                     --   many in parallel as there are CPUs.
```

```haskell
39
40    enableTimeStats      :: Bool,       -- ^ Enable RTS timing statistics?
41    ghcHeapSize          :: Maybe Int,  -- ^ The heap size to set.
42
43    maxRelevantBinds     :: Maybe Int,  -- ^ Maximum number of bindings from the type envt
44                                        --   to show in type error messages
45    maxValidHoleFits     :: Maybe Int,  -- ^ Maximum number of hole fits to show
46                                        --   in typed hole error messages
47    maxRefHoleFits       :: Maybe Int,  -- ^ Maximum number of refinement hole
48                                        --   fits to show in typed hole error
49                                        --   messages
50    refLevelHoleFits     :: Maybe Int,  -- ^ Maximum level of refinement for
51                                        --   refinement hole fits in typed hole
52                                        --   error messages
53    maxUncoveredPatterns :: Int,        -- ^ Maximum number of unmatched patterns to show
54                                        --   in non-exhaustiveness warnings
55    maxPmCheckModels     :: Int,        -- ^ Soft limit on the number of models
56                                        --   the pattern match checker checks
57                                        --   a pattern against. A safe guard
58                                        --   against exponential blow-up.
59    simplTickFactor      :: Int,        -- ^ Multiplier for simplifier ticks
60    dmdUnboxWidth        :: !Int,       -- ^ Whether DmdAnal should optimistically put an
61                                        --   Unboxed demand on returned products with at most
62                                        --   this number of fields
63    specConstrThreshold  :: Maybe Int,  -- ^ Threshold for SpecConstr
64    specConstrCount      :: Maybe Int,  -- ^ Max number of specialisations for any one function
65    specConstrRecursive  :: Int,        -- ^ Max number of specialisations for recursive types
66                                        --   Not optional; otherwise ForceSpecConstr can diverge.
67    binBlobThreshold     :: Maybe Word, -- ^ Binary literals (e.g. strings) whose size is above
68                                        --   this threshold will be dumped in a binary file
69                                        --   by the assembler code generator. 0 and Nothing disables
70                                        --   this feature. See 'GHC.StgToCmm.Config'.
71    liberateCaseThreshold :: Maybe Int, -- ^ Threshold for LiberateCase
72    floatLamArgs         :: Maybe Int,  -- ^ Arg count for lambda floating
73                                        --   See 'GHC.Core.Opt.Monad.FloatOutSwitches'
74
75    liftLamsRecArgs      :: Maybe Int,  -- ^ Maximum number of arguments after lambda lifting a
76                                        --   recursive function.
77    liftLamsNonRecArgs   :: Maybe Int,  -- ^ Maximum number of arguments after lambda lifting a
78                                        --   non-recursive function.
79    liftLamsKnown        :: Bool,       -- ^ Lambda lift even when this turns a known call
80                                        --   into an unknown call.
81
82    cmmProcAlignment     :: Maybe Int,  -- ^ Align Cmm functions at this boundary or use default.
83
84    historySize          :: Int,        -- ^ Simplification history size
85
86    importPaths          :: [FilePath],
87    mainModuleNameIs     :: ModuleName,
88    mainFunIs            :: Maybe String,
89    reductionDepth       :: IntWithInf, -- ^ Typechecker maximum stack depth
90    solverIterations     :: IntWithInf, -- ^ Number of iterations in the constraints solver
91                                        --   Typically only 1 is needed
92
93    homeUnitId_          :: UnitId,              -- ^ Target home unit-id
```

```haskell
    homeUnitInstanceOf_   :: Maybe UnitId,        -- ^ Id of the unit to instantiate
    homeUnitInstantiations_ :: [(ModuleName, Module)], -- ^ Module instantiations

    -- Note [Filepaths and Multiple Home Units]
    workingDirectory    :: Maybe FilePath,
    thisPackageName     :: Maybe String, -- ^ What the package is called, use with multiple home uni
    hiddenModules       :: Set.Set ModuleName,
    reexportedModules   :: Set.Set ModuleName,

    -- ways
    targetWays_         :: Ways,        -- ^ Target way flags from the command line

    -- For object splitting
    splitInfo           :: Maybe (String,Int),

    -- paths etc.
    objectDir           :: Maybe String,
    dylibInstallName    :: Maybe String,
    hiDir               :: Maybe String,
    hieDir              :: Maybe String,
    stubDir             :: Maybe String,
    dumpDir             :: Maybe String,

    objectSuf_          :: String,
    hcSuf               :: String,
    hiSuf_              :: String,
    hieSuf              :: String,

    dynObjectSuf_       :: String,
    dynHiSuf_           :: String,

    outputFile_         :: Maybe String,
    dynOutputFile_      :: Maybe String,
    outputHi            :: Maybe String,
    dynOutputHi         :: Maybe String,
    dynLibLoader        :: DynLibLoader,

    dynamicNow          :: !Bool, -- ^ Indicate if we are now generating dynamic output
                                -- because of -dynamic-too. This predicate is
                                -- used to query the appropriate fields
                                -- (outputFile/dynOutputFile, ways, etc.)

    -- | This defaults to 'non-module'. It can be set by
    -- 'GHC.Driver.Pipeline.setDumpPrefix' or 'ghc.GHCi.UI.runStmt' based on
    -- where its output is going.
    dumpPrefix          :: FilePath,

    -- | Override the 'dumpPrefix' set by 'GHC.Driver.Pipeline.setDumpPrefix'
    --   or 'ghc.GHCi.UI.runStmt'.
    --   Set by @-ddump-file-prefix@
    dumpPrefixForce     :: Maybe FilePath,

    ldInputs            :: [Option],

    includePaths        :: IncludeSpecs,
```

```haskell
149    libraryPaths        :: [String],
150    frameworkPaths      :: [String],   -- used on darwin only
151    cmdlineFrameworks   :: [String],   -- ditto
152
153    rtsOpts             :: Maybe String,
154    rtsOptsEnabled      :: RtsOptsEnabled,
155    rtsOptsSuggestions  :: Bool,
156
157    hpcDir              :: String,     -- ^ Path to store the .mix files
158
159    -- Plugins
160    pluginModNames      :: [ModuleName],
161      -- ^ the @-fplugin@ flags given on the command line, in *reverse*
162      -- order that they're specified on the command line.
163    pluginModNameOpts   :: [(ModuleName,String)],
164    frontendPluginOpts  :: [String],
165      -- ^ the @-ffrontend-opt@ flags given on the command line, in *reverse*
166      -- order that they're specified on the command line.
167
168    externalPluginSpecs :: [ExternalPluginSpec],
169      -- ^ External plugins loaded from shared libraries
170
171    -- For ghc -M
172    depMakefile         :: FilePath,
173    depIncludePkgDeps   :: Bool,
174    depIncludeCppDeps   :: Bool,
175    depExcludeMods      :: [ModuleName],
176    depSuffixes         :: [String],
177
178    -- Package flags
179    packageDBFlags      :: [PackageDBFlag],
180        -- ^ The @-package-db@ flags given on the command line, In
181        -- *reverse* order that they're specified on the command line.
182        -- This is intended to be applied with the list of "initial"
183        -- package databases derived from @GHC_PACKAGE_PATH@; see
184        -- 'getUnitDbRefs'.
185
186    ignorePackageFlags  :: [IgnorePackageFlag],
187        -- ^ The @-ignore-package@ flags from the command line.
188        -- In *reverse* order that they're specified on the command line.
189    packageFlags        :: [PackageFlag],
190        -- ^ The @-package@ and @-hide-package@ flags from the command-line.
191        -- In *reverse* order that they're specified on the command line.
192    pluginPackageFlags  :: [PackageFlag],
193        -- ^ The @-plugin-package-id@ flags from command line.
194        -- In *reverse* order that they're specified on the command line.
195    trustFlags          :: [TrustFlag],
196        -- ^ The @-trust@ and @-distrust@ flags.
197        -- In *reverse* order that they're specified on the command line.
198    packageEnv          :: Maybe FilePath,
199        -- ^ Filepath to the package environment file (if overriding default)
200
201
202    -- hsc dynamic flags
203    dumpFlags           :: EnumSet DumpFlag,
```

```
204    generalFlags        :: EnumSet GeneralFlag,
205    warningFlags        :: EnumSet WarningFlag,
206    fatalWarningFlags   :: EnumSet WarningFlag,
207    -- Don't change this without updating extensionFlags:
208    language            :: Maybe Language,
209    -- | Safe Haskell mode
210    safeHaskell         :: SafeHaskellMode,
211    safeInfer           :: Bool,
212    safeInferred        :: Bool,
213    -- We store the location of where some extension and flags were turned on so
214    -- we can produce accurate error messages when Safe Haskell fails due to
215    -- them.
216    thOnLoc             :: SrcSpan,
217    newDerivOnLoc       :: SrcSpan,
218    deriveViaOnLoc      :: SrcSpan,
219    overlapInstLoc      :: SrcSpan,
220    incoherentOnLoc     :: SrcSpan,
221    pkgTrustOnLoc       :: SrcSpan,
222    warnSafeOnLoc       :: SrcSpan,
223    warnUnsafeOnLoc     :: SrcSpan,
224    trustworthyOnLoc    :: SrcSpan,
225    -- Don't change this without updating extensionFlags:
226    -- Here we collect the settings of the language extensions
227    -- from the command line, the ghci config file and
228    -- from interactive :set / :seti commands.
229    extensions          :: [OnOff LangExt.Extension],
230    -- extensionFlags should always be equal to
231    --     flattenExtensionFlags language extensions
232    -- LangExt.Extension is defined in libraries/ghc-boot so that it can be used
233    -- by template-haskell
234    extensionFlags      :: EnumSet LangExt.Extension,
235
236    -- | Unfolding control
237    -- See Note [Discounts and thresholds] in GHC.Core.Unfold
238    unfoldingOpts       :: !UnfoldingOpts,
239
240    maxWorkerArgs       :: Int,
241
242    ghciHistSize        :: Int,
243
244    flushOut            :: FlushOut,
245
246    ghcVersionFile      :: Maybe FilePath,
247    haddockOptions      :: Maybe String,
248
249    -- | GHCi scripts specified by -ghci-script, in reverse order
250    ghciScripts         :: [String],
251
252    -- Output style options
253    pprUserLength       :: Int,
254    pprCols             :: Int,
255
256    useUnicode          :: Bool,
257    useColor            :: OverridingBool,
258    canUseColor         :: Bool,
```

```haskell
259    colScheme          :: Col.Scheme,
260
261    -- | what kind of {-# SCC #-} to add automatically
262    profAuto           :: ProfAuto,
263    callerCcFilters    :: [CallerCcFilter],
264
265    interactivePrint   :: Maybe String,
266
267    -- | Machine dependent flags (-m\<blah> stuff)
268    sseVersion         :: Maybe SseVersion,
269    bmiVersion         :: Maybe BmiVersion,
270    avx                :: Bool,
271    avx2               :: Bool,
272    avx512cd           :: Bool, -- Enable AVX-512 Conflict Detection Instructions.
273    avx512er           :: Bool, -- Enable AVX-512 Exponential and Reciprocal Instructions.
274    avx512f            :: Bool, -- Enable AVX-512 instructions.
275    avx512pf           :: Bool, -- Enable AVX-512 PreFetch Instructions.
276
277    -- | Run-time linker information (what options we need, etc.)
278    rtldInfo           :: IORef (Maybe LinkerInfo),
279
280    -- | Run-time C compiler information
281    rtccInfo           :: IORef (Maybe CompilerInfo),
282
283    -- | Run-time assembler information
284    rtasmInfo          :: IORef (Maybe CompilerInfo),
285
286    -- Constants used to control the amount of optimization done.
287
288    -- | Max size, in bytes, of inline array allocations.
289    maxInlineAllocSize  :: Int,
290
291    -- | Only inline memcpy if it generates no more than this many
292    -- pseudo (roughly: Cmm) instructions.
293    maxInlineMemcpyInsns :: Int,
294
295    -- | Only inline memset if it generates no more than this many
296    -- pseudo (roughly: Cmm) instructions.
297    maxInlineMemsetInsns :: Int,
298
299    -- | Reverse the order of error messages in GHC/GHCi
300    reverseErrors      :: Bool,
301
302    -- | Limit the maximum number of errors to show
303    maxErrors          :: Maybe Int,
304
305    -- | Unique supply configuration for testing build determinism
306    initialUnique      :: Word,
307    uniqueIncrement    :: Int,
308      -- 'Int' because it can be used to test uniques in decreasing order.
309
310    -- | Temporary: CFG Edge weights for fast iterations
311    cfgWeights         :: Weights
312 }
```

Appendix B

GHC DynFlags default value

Sourced from [24].

```
1   -- | The normal 'DynFlags'. Note that they are not suitable for use in this form
2   -- and must be fully initialized by 'GHC.runGhc' first.
3   defaultDynFlags :: Settings -> DynFlags
4   defaultDynFlags mySettings =
5   -- See Note [Updating flag description in the User's Guide]
6       DynFlags {
7           ghcMode             = CompManager,
8           ghcLink             = LinkBinary,
9           backend             = platformDefaultBackend (sTargetPlatform mySettings),
10          verbosity           = 0,
11          debugLevel          = 0,
12          simplPhases         = 2,
13          maxSimplIterations  = 4,
14          ruleCheck           = Nothing,
15          binBlobThreshold    = Just 500000, -- 500K is a good default (see #16190)
16          maxRelevantBinds    = Just 6,
17          maxValidHoleFits = Just 6,
18          maxRefHoleFits   = Just 6,
19          refLevelHoleFits = Nothing,
20          maxUncoveredPatterns = 4,
21          maxPmCheckModels    = 30,
22          simplTickFactor     = 100,
23          dmdUnboxWidth       = 3,      -- Default: Assume an unboxed demand on function bodies retu
24          specConstrThreshold = Just 2000,
25          specConstrCount     = Just 3,
26          specConstrRecursive = 3,
27          liberateCaseThreshold = Just 2000,
28          floatLamArgs        = Just 0, -- Default: float only if no fvs
29          liftLamsRecArgs     = Just 5, -- Default: the number of available argument hardware regis
30          liftLamsNonRecArgs  = Just 5, -- Default: the number of available argument hardware regis
31          liftLamsKnown       = False, -- Default: don't turn known calls into unknown ones
32          cmmProcAlignment    = Nothing,
33
34          historySize         = 20,
35          strictnessBefore    = [],
36
37          parMakeCount        = Just 1,
38
```

41

```
39          enableTimeStats       = False,
40          ghcHeapSize           = Nothing,
41
42          importPaths           = ["."],
43          mainModuleNameIs      = mAIN_NAME,
44          mainFunIs             = Nothing,
45          reductionDepth        = treatZeroAsInf mAX_REDUCTION_DEPTH,
46          solverIterations      = treatZeroAsInf mAX_SOLVER_ITERATIONS,
47
48          homeUnitId_           = mainUnitId,
49          homeUnitInstanceOf_   = Nothing,
50          homeUnitInstantiations_ = [],
51
52          workingDirectory      = Nothing,
53          thisPackageName       = Nothing,
54          hiddenModules         = Set.empty,
55          reexportedModules     = Set.empty,
56
57          objectDir             = Nothing,
58          dylibInstallName      = Nothing,
59          hiDir                 = Nothing,
60          hieDir                = Nothing,
61          stubDir               = Nothing,
62          dumpDir               = Nothing,
63
64          objectSuf_            = phaseInputExt StopLn,
65          hcSuf                 = phaseInputExt HCc,
66          hiSuf_                = "hi",
67          hieSuf                = "hie",
68
69          dynObjectSuf_         = "dyn_" ++ phaseInputExt StopLn,
70          dynHiSuf_             = "dyn_hi",
71          dynamicNow            = False,
72
73          pluginModNames        = [],
74          pluginModNameOpts     = [],
75          frontendPluginOpts    = [],
76
77          externalPluginSpecs   = [],
78
79          outputFile_           = Nothing,
80          dynOutputFile_        = Nothing,
81          outputHi              = Nothing,
82          dynOutputHi           = Nothing,
83          dynLibLoader          = SystemDependent,
84          dumpPrefix            = "non-module.",
85          dumpPrefixForce       = Nothing,
86          ldInputs              = [],
87          includePaths          = IncludeSpecs [] [] [],
88          libraryPaths          = [],
89          frameworkPaths        = [],
90          cmdlineFrameworks     = [],
91          rtsOpts               = Nothing,
92          rtsOptsEnabled        = RtsOptsSafeOnly,
93          rtsOptsSuggestions    = True,
```

```
94
95          hpcDir                = ".hpc",
96
97          packageDBFlags      = [],
98          packageFlags        = [],
99          pluginPackageFlags  = [],
100         ignorePackageFlags  = [],
101         trustFlags          = [],
102         packageEnv          = Nothing,
103         targetWays_         = Set.empty,
104         splitInfo           = Nothing,
105
106         ghcNameVersion = sGhcNameVersion mySettings,
107         fileSettings = sFileSettings mySettings,
108         toolSettings = sToolSettings mySettings,
109         targetPlatform = sTargetPlatform mySettings,
110         platformMisc = sPlatformMisc mySettings,
111         rawSettings = sRawSettings mySettings,
112
113         tmpDir                = panic "defaultDynFlags: uninitialized tmpDir",
114
115         llvmOptLevel          = 0,
116
117         -- ghc -M values
118         depMakefile     = "Makefile",
119         depIncludePkgDeps = False,
120         depIncludeCppDeps = False,
121         depExcludeMods  = [],
122         depSuffixes     = [],
123         -- end of ghc -M values
124         ghcVersionFile = Nothing,
125         haddockOptions = Nothing,
126         dumpFlags = EnumSet.empty,
127         generalFlags = EnumSet.fromList (defaultFlags mySettings),
128         warningFlags = EnumSet.fromList standardWarnings,
129         fatalWarningFlags = EnumSet.empty,
130         ghciScripts = [],
131         language = Nothing,
132         safeHaskell = Sf_None,
133         safeInfer  = True,
134         safeInferred = True,
135         thOnLoc = noSrcSpan,
136         newDerivOnLoc = noSrcSpan,
137         deriveViaOnLoc = noSrcSpan,
138         overlapInstLoc = noSrcSpan,
139         incoherentOnLoc = noSrcSpan,
140         pkgTrustOnLoc = noSrcSpan,
141         warnSafeOnLoc = noSrcSpan,
142         warnUnsafeOnLoc = noSrcSpan,
143         trustworthyOnLoc = noSrcSpan,
144         extensions = [],
145         extensionFlags = flattenExtensionFlags Nothing [],
146
147         unfoldingOpts = defaultUnfoldingOpts,
148         maxWorkerArgs = 10,
```

```
149
150        ghciHistSize = 50, -- keep a log of length 50 by default
151
152        flushOut = defaultFlushOut,
153        pprUserLength = 5,
154        pprCols = 100,
155        useUnicode = False,
156        useColor = Auto,
157        canUseColor = False,
158        colScheme = Col.defaultScheme,
159        profAuto = NoProfAuto,
160        callerCcFilters = [],
161        interactivePrint = Nothing,
162        sseVersion = Nothing,
163        bmiVersion = Nothing,
164        avx = False,
165        avx2 = False,
166        avx512cd = False,
167        avx512er = False,
168        avx512f = False,
169        avx512pf = False,
170        rtldInfo = panic "defaultDynFlags: no rtldInfo",
171        rtccInfo = panic "defaultDynFlags: no rtccInfo",
172        rtasmInfo = panic "defaultDynFlags: no rtasmInfo",
173
174        maxInlineAllocSize = 128,
175        maxInlineMemcpyInsns = 32,
176        maxInlineMemsetInsns = 32,
177
178        initialUnique = 0,
179        uniqueIncrement = 1,
180
181        reverseErrors = False,
182        maxErrors    = Nothing,
183        cfgWeights   = defaultWeights
184    }
```

Bibliography

[1] ANDY GILL, S. M. Happy, a parser generator for haskell. https://www.haskell.org/happy/. Accessed: 2022-11-20.

[2] BRAINERD, W. Fortran 77. *Commun. ACM 21*, 10 (oct 1978).

[3] DAHL, O.-J. *SIMULA 67 Common Base Language, (Norwegian Computing Center. Publication)*. 1968.

[4] DIEHL, S. Dive into ghc: Intermediate forms. https://www.stephendiehl.com/posts/ghc_02.html, 2016. Accessed: 2022-11-12.

[5] ERICSSON AB. Erlang programming examples user's guide - records. https://www.erlang.org/doc/programming_examples/records.html. Accessed: 2022-12-03.

[6] GOLDBERG, A., AND ROBSON, D. *Smalltalk-80: The Language and Its Implementation*. Addison-Wesley Longman Publishing Co., Inc., USA, 1983.

[7] HUDAK, P., PEYTON JONES, S., WADLER, P., BOUTEL, B., FAIRBAIRN, J., FASEL, J., GUZMÁN, M., HAMMOND, K., HUGHES, J., JOHNSSON, T., KIEBURTZ, R., NIKHIL, R., PARTAIN, W., AND PETERSON, J. Report on the programming language haskell: a non-strict, purely functional language version 1.2. *SIGPLAN Notices 27* (01 1992), 1–.

[8] ISO. *ISO/IEC 14882:1998: Programming languages — C++*. pub-ISO, pub-ISO:adr, Sept. 1998. Available in electronic form for online purchase at http://webstore.ansi.org/ and http://www.cssinfo.com/.

[9] ISO. *ISO/IEC 14882:2011 Information technology — Programming languages — C++*, third ed. pub-ISO, pub-ISO:adr, Sept. 2011.

[10] JONES, S. P., HAMMOND, K., PARTAIN, W., WADLER, P., HALL, C. B., AND JONES, S. L. P. The glasgow haskell compiler: a technical overview.

[11] NAJD, S., AND JONES, S. L. P. Trees that grow. *ArXiv abs/1610.04799* (2017).

[12] PARTAIN, W., SANTOS, A., AND PEYTON JONES, S. Let-floating: moving bindings to give faster programs, May 1996. ACM SIGPLAN International Conference on Functional Programming (ICFP'96).

[13] SALTZER, J., AND KAASHOEK, F. *Principles of Computer System Design: An Introduction*. Morgan Kaufmann, 2009.

[14] SEBESTA, R. *Concepts of Programming Languages*. Benjamin/Cummings series in computer science. Addison-Wesley Publishing Company, 1996.

[15] SIMON MARLOW. Haddock docs - documenting a top-level declaration. https://haskell-haddock.readthedocs.io/en/v2.20/markup.html. Accessed: 2022-12-04.

[16] STATISTA. Most used programming languages among developers worldwide as of 2022. https://www.statista.com/statistics/793628/worldwide-developer-survey-most-used-languages/. Accessed: 2022-11-25.

[17] STEELE, G. L. *Common LISP: The Language.* Digital Press, USA, 1984.

[18] THE GLASGOW HASKELL TEAM. Ghc 9.4.2 user's guide - 6.13 template haskell. https://downloads.haskell.org/~ghc/9.4.2/docs/users_guide/exts/template_haskell.html. Accessed: 2022-12-06.

[19] THE GLASGOW HASKELL TEAM. Ghc wiki commentary: The core language. https://gitlab.haskell.org/ghc/ghc/-/wikis/commentary/compiler/core-syn-type?version_id=41a3eea9a2e430e1f61fb00bfa3f9d4210ada149. Accessed: 2022-11-21.

[20] THE GLASGOW HASKELL TEAM. Ghc wiki commentary: The parser. https://gitlab.haskell.org/ghc/ghc/-/wikis/commentary/compiler/parser?version_id=f2e50659ba9fd4cd8f4231d498a16fe67dcb0013. Accessed: 2022-11-20.

[21] THE GLASGOW HASKELL TEAM. Ghc wiki commentary: The renamer. https://gitlab.haskell.org/ghc/ghc/-/wikis/commentary/compiler/renamer?version_id=93c96a9daa6b08a02b51c151e5e1f7b19194fa79. Accessed: 2022-11-21.

[22] THE GLASGOW HASKELL TEAM. Ghc wiki commentary: The typechecker. https://gitlab.haskell.org/ghc/ghc/-/wikis/commentary/compiler/type-checker?version_id=f95129964a2c0a5ae276bc012bea63b05a415b76. Accessed: 2022-11-21.

[23] THE GLASGOW HASKELL TEAM. ghc/compiler/ghc/core.hs - expr type definition. https://gitlab.haskell.org/ghc/ghc/-/blob/1a767fa359d22ca7637af41e29434e76487c3f21/compiler/GHC/Core.hs#L249-261. Accessed: 2022-12-06.

[24] THE GLASGOW HASKELL TEAM. ghc/compiler/ghc/driver/session.hs - dynflags default value. https://gitlab.haskell.org/ghc/ghc/-/blob/1a767fa359d22ca7637af41e29434e76487c3f21/compiler/GHC/Driver/Session.hs#L1111-1292. Accessed: 2022-12-06.

[25] THE GLASGOW HASKELL TEAM. ghc/compiler/ghc/driver/session.hs - dynflags type definition. https://gitlab.haskell.org/ghc/ghc/-/blob/1a767fa359d22ca7637af41e29434e76487c3f21/compiler/GHC/Driver/Session.hs#L429-740. Accessed: 2022-12-06.

[26] THE GLASGOW HASKELL TEAM. ghc/compiler/ghc/types/fieldlabel.hs - fieldlabel type definition. https://gitlab.haskell.org/minimario/ghc/-/blob/5b41353355022c1247e0516d541b7f7fb49f0e29/compiler/GHC/Types/FieldLabel.hs#L105-117. Accessed: 2022-12-06.

[27] THE GLASGOW HASKELL TEAM. ghc/compiler/language/haskell/syntax/type.hs - condeclfield type definition. https://gitlab.haskell.org/minimario/ghc/-/blob/5b41353355022c1247e0516d541b7f7fb49f0e29/compiler/Language/Haskell/Syntax/Type.hs#L1043-1050. Accessed: 2022-12-06.

[28] van Wijngaarden, A., Mailloux, B., Peck, J., Koster, C., Lindsey, C., Sint-zoff, M., Meertens, L., and Fisker, R. *Revised Report on the Algorithmic Language Algol 68.* Springer Berlin Heidelberg, 2012.

[29] Wikibooks. Haskell/more on datatypes. `https://en.wikibooks.org/w/index.php?title=Haskell/More_on_datatypes&oldid=3676091`. Accessed: 2022-11-23.